

Alibaba Group

June Quarter 2026 Results

Conference Call

Transcript

Thursday, 20 August 2026

Introduction

Lydia Liu

Head of Investor Relations, Alibaba Group

(Original)

Thank you. Good day everyone, and welcome to Alibaba Group's June Quarter 2026 earnings conference call.

On the call with me are Joe Tsai, Chairman; Eddie Wu, Chief Executive Officer; Toby Xu, Chief Financial Officer; Jiang Fan, Chief Executive Officer of Alibaba E-commerce Business Group.

I would like to remind you that this call is also being webcast on our corporate website. A replay of the call will be available on our website later today.

Now I will quickly cover the safe harbor. Today's discussions may contain forward-looking statements, particularly statements about our business prospects and expected financial results, that are subject to risks and uncertainties, which could cause actual results to differ materially from those contained in the forward-looking statements. Please refer to the safe harbor statements that appear in our press release and investor presentation provided today. Please note that certain financial measures that we use on this call are expressed on a non-GAAP basis. Our GAAP results and reconciliations of GAAP to non-GAAP measures can be found in our earnings press release and investor presentation provided today.

Unless otherwise stated, growth rate of all stated metrics mentioned during this call refers to year-over-year growth versus the same quarter last year.

With that, I would like to turn the call over to Eddie.

Group Business Update

Eddie Wu

CEO, Alibaba Group

(Translation)

Dear investors, welcome to Alibaba Group's Earnings Call for the First Quarter of Fiscal Year 2027.

Over the past quarter, Alibaba's strategic AI investments have translated into robust results with total Group revenue growing 9% year-over-year. AI commercialization has also accelerated across the board: Alibaba Cloud's revenue from external customers grew a robust 45% and adjusted EBITA increased 133% year-over-year, continuing to deliver on our commitment to accelerate growth. Revenue from AI-related products maintained triple-digit growth for the twelfth consecutive quarter, with the annualized revenue run rate surpassing RMB49.5 billion (around USD7.3 billion). This is the core engine of Alibaba Cloud's accelerating growth.

I'll now walk you through four key areas: AI and cloud commercialization, full-stack AI, the AI application ecosystem, and our consumption businesses.

First, AI and cloud commercialization accelerated across the board and is expected to sustain high growth going forward.

This quarter, Alibaba Cloud's revenue from external customers growth accelerated to 45%, a 22-quarter high, while adjusted EBITA margin of 11.6%. Notably, this 45% growth was broad-based, driven by compute, storage, MaaS, and AI applications. We proactively scaled back our some low-margin network business, continuing to enhance the quality of our growth.

This quarter, the annualized revenue run rate of AI-related products exceeded RMB49.5 billion, while AI-related products' share of Alibaba Cloud's revenue from external customers rose to 35%. AI-related products generate significantly higher gross margins than the cloud portfolio average. Our recurring revenue from AI-related products spans multiple layers, including AI compute, model services, and AI applications. This multi-layered mix of AI revenue sources and monetization models means that growth in customer demand at any layer — compute, models, or applications — will translate directly into commercial opportunities for us. This structural advantage will underpin sustained rapid growth in recurring AI-related product revenue going forward.

The explosive growth of AI agents is directly driving demand for tokens and GPU compute, while also significantly boosting demand for our traditional cloud products, including CPU compute, storage, databases, and networking. Alibaba Cloud is undergoing a comprehensive upgrade to

an Agentic Cloud. Based on the latest data for August, the ARR of our AI model and application services, including MaaS, has surpassed RMB16 billion.

Based on current market feedback and contracted orders, compute demand will continue to outstrip supply. As we continue to ramp up our supply capacity, our AI and cloud revenue growth will accelerate further in the coming quarters, alongside continued improvement in profitability.

Second, our full-stack AI continue to strengthen, marked by the scaled commercialization of proprietary chips, faster model iteration, and a thriving open-source ecosystem.

This quarter, the deep synergy between T-Head's proprietary chips and our proprietary models further enhanced our AI commercialization efficiency.

T-Head has established a full-stack portfolio of proprietary chips spanning GPUs, CPUs, and networking chips. As of early August, Zhenwu chips had served more than 650 customers. The supernode instance powered by T-Head's next-generation Zhenwu M890 AI chip, backed entirely by domestic supply chain, recently launched on Alibaba Cloud for commercial sale at scale. Volumes will continue to ramp up in the second half of the year to meet strong customer demand. Alibaba Cloud's Zhenwu M890 Supernode can efficiently run inference workloads for foundation models with more than 2 trillion parameters. Both Kimi K3 and Qwen 3.8-Max are already providing MaaS services to external customers through this supernode instance.

At the data center layer, Alibaba Cloud has reduced the delivery time for large AI Data Centers (AIDCs) to 100 days, a world-leading pace that will significantly speed up our global compute infrastructure build-out.

At the model layer, our model releases have accelerated across the board over the past month, with major iterations across our large language, image, speech, video, and music models, all with performance ranking in the top tier globally.

Last week, we opened the model weights of Qwen 3.8-Max with 2.4-trillion-parameters and the Qwen 3.8-27B model series. To date, the Qwen model series has been downloaded more than 3 billion times globally, with over 300,000 derivative models built on it. A thriving open-source model ecosystem will help drive greater demand for cloud services, creating a virtuous commercial cycle.

Third, we are developing AI-native applications for both enterprise and consumer use cases, driving rapid growth in token consumption.

On the enterprise side, we launched QwenWork, a new AI productivity product designed for enterprise workplace use cases, delivering foundation model and agent capabilities at scale. We believe that productivity agents will become another engine of ARR growth.

On the consumer side, the Qwen app has continued to grow its user base steadily while adding a broader range of paid value-added services.

With Alibaba Token Hub (ATH) and Alibaba Cloud working in close coordination, the complete "compute–model–token–application–revenue" commercial chain is operating efficiently.

Fourth, our e-commerce businesses remained stable overall this quarter. In quick commerce, we substantially narrowed losses while maintaining high growth of 45% in business scale, with UE (unit economics) continuing to improve quarter-over-quarter.

Having crossed the AI commercialization inflection point last quarter, we saw growth accelerate and margins expand this quarter. Our AI business is becoming increasingly self-funding and better able to support continued investment, giving us even greater confidence to keep investing. Looking ahead, AI has become Alibaba's most certain growth engine. We will maintain our strategic resolve and use our full-stack AI to drive long-term growth.

I'll now hand over to Toby to walk you through our financial results. Thank you.

Financial Highlights

Toby Xu

CFO, Alibaba Group

(Original)

Thank you, Eddie.

Our strategic priorities in AI+Cloud and consumption businesses, backed by disciplined investments, delivered strong results this quarter.

Cloud segment revenue growth further accelerated to 45%, with its EBITA margin sequentially rising to 12%. AI-related product revenue continued to drive this momentum, marking the twelfth consecutive quarter of triple-digit growth and accounting for 35% of external cloud revenue. The strong performance demonstrates growing customer adoption of our full-stack AI capabilities spanning AI agents, models, cloud infrastructure and proprietary chips, as well as our enhanced scale efficiencies and robust pricing power in a supply-constrained market.

On consumption, Taobao Instant Commerce continued to improve its UE, while maintaining market share. Overall e-commerce EBITA remained relatively stable year-over-year.

To realize synergies across our commerce platforms and strengthen our full-stack AI capabilities, we have implemented strategic re-alignment of certain businesses in our financial reporting. Starting from this quarter, our segment reporting will present the following: (1) Alibaba E-commerce Group, (2) AI Cloud and Compute Services, (3) AI Labs and Applications, and (4) All Others.

Now let's look at the financial results for this quarter:

Total revenue increased 9% year-over-year to RMB269 billion, driven by the strong momentum in Cloud business and quick commerce.

Total adjusted EBITA decreased 30% to RMB27.3 billion, primarily attributable to the investment in technology, partly offset by the improved operating results in our Cloud business, as well as enhanced operating efficiencies across various businesses.

Our GAAP net income was RMB10.4 billion, a decrease of 75%, primarily due to the decrease in income from operations, and decrease in net gains from disposal of investments and mark-to-market changes of our equity investments.

Operating cash flow this quarter increased by 11% to RMB22.9 billion, compared to RMB20.7 billion in the same quarter last year.

Free cash flow was an outflow of RMB44.7 billion, compared to an outflow of RMB18.8 billion in the same quarter last year. The decrease was mainly attributed to the investment in cloud infrastructure.

CapEx was RMB67.7 billion this quarter, reflecting our continued investments in AI infrastructure to meet strong and growing customer demand. The significant year-over-year increase is due to several reasons, including fluctuations in procurement cycles, increase in CPU-compute capacity driven by anticipated growing customer adoption of Al agents, and higher pricing of a broad range of chip components.

As of June 30, 2026, we held approximately USD30.7 billion in net cash. Excluding debt with maturities beyond five years, our net cash position stands at approximately USD46.5 billion. This balance sheet strength gives us confidence to invest for robust growth.

Our Al+Cloud investment has a clear path to attractive ROIC. Our servers equipped with chips typically reach break-even within three years. With a five-year useful life, we expect them to generate positive free cash flow at least in the two years following break-even.

For the quarter ended June 30, 2026, we repurchased shares for an aggregate consideration of USD162 million.

We remain committed to maximizing long-term shareholder returns through disciplined capital allocation across investments for Al+Cloud business growth, share buybacks and dividends. We will adjust our priorities as market conditions and strategic needs evolve.

Now let's first look at our E-commerce businesses.

The new Alibaba E-commerce Group reflects our strategic focus on unlocking significant synergies across our domestic and cross-border e-commerce businesses.

Starting from this quarter, we will present Alibaba E-commerce Group's revenue as the following: (1) China E-commerce, (2) China Quick Commerce, (3) International E-commerce, and (4) Global Wholesale.

Revenue for Alibaba E-commerce Group was RMB205.9 billion, an increase of 4%.

Customer management revenue decreased by 7%. Excluding the contra revenue impact from the new business development program, customer management revenue would have grown by 1% year-over-year.

Revenue from China Quick Commerce business was RMB53.3 billion, an increase of 45%, driven by Freshippo and Taobao Instant Commerce.

Alibaba E-commerce Group's adjusted EBITA remained relatively stable year-over-year at RMB39.7 billion, underscoring our cost discipline against the backdrop of increased investment in user experiences and technology.

Taobao Instant Commerce continued to improve its UE quarter-over-quarter while maintaining market share, driven by higher average order value and enhanced fulfillment logistics efficiency.

In addition, AliExpress achieved operating profit this quarter.

We aim to maintain steady profit in our conventional e-commerce business, while continuing to drive profitability improvement in our quick commerce business.

Now let's review the business update and results of AI Cloud and Compute Services, which comprises the Cloud Intelligence Group and T-Head.

The year-over-year growth of total revenue and revenue from external customers both accelerated to 45%. Revenue from Alibaba Cloud also accelerated, growing 45% year-over-year. We are confident the growth rate will further accelerate in the coming quarters.

This quarter's AI-related product revenue was RMB12.4 billion, implying an annual revenue run rate of RMB49.5 billion. It delivered the twelfth consecutive quarter of triple-digit growth and accounted for 35% of external cloud revenue.

The adjusted EBITA margin expanded to 12%, driven by improved economies of scale and stronger pricing power for AI-related products amid tight market supply.

We expect adjusted EBITA margin to further expand steadily in the coming quarters. By improving resource utilization, optimizing model portfolio, and innovating new scenarios, we are accelerating the growth of AI+Cloud business and driving greater benefits of scale.

AI Labs and Applications comprises AI model labs, Qwen Consumer Business Group, and QwenWork.

Its adjusted EBITA was a loss of RMB13.9 billion, primarily due to our increased investment in AI capabilities and higher inference costs related to Qwen app. The loss significantly narrowed quarter-over-quarter due to the reduction in marketing expenses for Qwen app. We expect the segment loss to narrow over the coming quarters, driven by improving efficiency in both model training and marketing spend on Qwen app.

We have launched our frontier language, coding, video, audio, image and music models, all delivering top-tier performance.

250 million users have had their first AI-driven shopping experience through Qwen app's agentic features across an expanding range of e-commerce and other services since the launch of Qwen app.

All Others segment revenue remained stable at RMB28.8 billion.

All Others adjusted EBITA was a loss of RMB3.3 billion, primarily due to our increased investment in technology.

AI has progressed from incubation to commercialization at scale. As we expand our market share, strengthen AI leadership and improve operating efficiency, we are gaining greater

strategic and financial flexibility to make disciplined and sustained investments in both full-stack AI capabilities and consumption opportunities, driving secular growth and greater value for our shareholders.

Thank you, that's the end of our prepared remarks. We can open up for Q&A.

Q&A

Lydia Liu:

(Original)

Thank you, Toby. We will now begin the Q&A session. You're welcome to ask questions in Chinese or English. A third-party translator will provide consecutive interpretation. In the case of any discrepancy, our management statements in the original language will prevail.

Operator, please start the Q&A session. Thank you.

Alicia Yap (Citigroup):

(Original)

Congrats on your solid cloud performance.

Could management please comment on the reasons and the drivers for the significant increase in the CapEx this quarter? And also, what is the expected CapEx trend for the coming quarters? And are there any updates to the existing 3-year CapEx budget that you have of this RMB 380 billion that you mentioned before?

And also, we would appreciate if management can also provide a breakdown of the CapEx allocation across the different services, like the training costs and all that? And then also, what is management's expected return on the invested capital for these investments?

Eddie Wu:

(Translation)

Thank you for your questions. These are important questions, and I would like to take this opportunity to explain our AI business model in more detail, as well as our expectations for future CapEx investment.

In February last year, we announced our three-year RMB380 billion CapEx plan, and as of the end of the June quarter we had invested over RMB190 billion cumulatively, with progress in line with expectations. CapEx of RMB67.7 billion this quarter was indeed a bit on the high side. This was largely due to fluctuations in equipment deliveries, as hardware deliveries follow their own cycles and are not evenly distributed across quarters. In addition, we increased CPU

procurement this quarter, as we are seeing rapidly growing demand for CPUs in the era of AI agents. Of course, rising prices for chip components were also a contributing factor.

So one should not take this quarter's RMB67.7 billion in CapEx, multiply it by four and assume that investment will proceed at a uniform quarterly pace throughout the year. That said, we do intend to remain proactive in our investment and infrastructure build-out this year.

I would also like to take this opportunity to explain the business model of our full-stack AI platform services. It is fundamentally an asset-heavy business model. All forms of AI monetization we see today — including AI software subscriptions, foundation model API inference services, GPU-based IaaS services, and software services for AI training and inference — ultimately depend on AI compute infrastructure. Only by building out AI compute centers can we capture rapid market-share growth. This is therefore clearly an asset-heavy business model, and rapid growth under such a model requires front-loaded CapEx.

That is why we entered a very proactive cycle of investment in and build-out of hardware infrastructure beginning in 2025. To achieve rapid growth, we must invest more in AI compute centers in the early years. This is inherent in our business model: CapEx must be deployed upfront before the subsequent business growth can materialize.

Let me also explain why returns on AI compute-related CapEx are highly certain at this stage. The industry consensus is that the shortage of AI compute will not change materially before 2030. Against this industry backdrop, we see very high certainty in the returns from our AI compute CapEx investments. Looking at the numbers, at the current average gross margin level of our AI products, we can recoup our CapEx investment within three years. And as the gross margins of our AI products continue to improve, we expect the payback period to shorten further, for example toward 2.5 years.

Based on our current payback cycle, after the initial investment is recouped in three years, these AI compute assets will continue to generate positive cash flow for a very long time. One of the most direct examples we see today is that, in our data centers, A100 GPUs purchased in 2020 and V100 GPUs purchased in 2018 continue to be used by customers at close to full capacity. The actual useful life of AI compute assets is substantially longer than their theoretical depreciation periods.

At the same time, we have three key levers to further increase the gross margins of our AI products and our returns on investment.

First, we can continue to enhance the gross margin of our AI products. This includes our ongoing investment in advancing the development of SOTA models and expanding our higher-margin MaaS business. By offering a more comprehensive product portfolio, we can also strengthen the competitiveness of our AI IaaS and AI-related software products, driving higher gross margins. AI product gross margin continues to improve, and this quarter it drove a 4.4 percentage-point increase in the overall margin of Alibaba Cloud, bringing it to 11.6%. This provides initial validation.

Second, we can leverage our proprietary chips. This is perhaps the most important initiative in this process. T-Head's proprietary chip roadmap spans GPUs, CPUs, and networking chips — the core chip portfolio for next-generation AI data centers. Chips and storage represent the largest cost components in AI data centers, and proprietary chips are a key long-term priority for us. As T-Head's production capacity continues to ramp up, proprietary chips will account for a growing share of our data center deployments, displacing externally procured chips. As we know, in today's supply-constrained compute market, commercial chips themselves command very high gross margins. A significantly higher share of proprietary chips will therefore substantially strengthen our product competitiveness and increase our gross margins.

Third, and also very importantly, we have several commercial arrangements to improve cash flow efficiency. For example, we can jointly build AI compute centers with partners, and some of our AI compute contracts have already begun to incorporate customer prepayments. These commercial approaches can also materially improve our returns on investment.

Together, these three levers can continue to shorten the payback period for AI CapEx towards 2.5 years and potentially close to 2 years.

Let me walk you through a simple framework. At current AI product gross margins and assuming a three-year payback period on CapEx, we could theoretically generate positive cash flow if we kept our growth rate at or below 33%. However, this is not our strategic choice at this stage. The AI industry remains at a very early stage, and our strategic choice is to remain firmly committed to proactive CapEx investment and build-out in order to drive rapid business growth.

As our product gross margins increase, the share of proprietary chips in our deployments rises, and the payback period shortens to 2.5 years or less, we believe we will be able to maintain positive cash flow while pursuing growth of more than 40%. That is our long-term direction.

Charlene Liu (HSBC):

(Original)

First, can we get an update on the latest developments in quick commerce and under the reclassification of multiple business lines, which are regrouped under the Alibaba E-commerce Group. Can you talk about the future strategic focuses of these lines of businesses?

Jiang Fan:

(Translation)

Let me start with an overview. In the new fiscal year, we have realigned our e-commerce business segments. Going forward, we will provide regular updates on four core areas: China E-commerce, Quick Commerce, International E-commerce, and Global Wholesale. Today, I would like to briefly share our progress and thinking in each of these areas.

First, China E-commerce. In the short term, domestic e-commerce still faces macroeconomic challenges. Taking a long-term view, we are focusing on core supply, while aiming to leverage AI to enhance the shopping experience and improve overall operational efficiency.

Regarding supply: since last year, Taobao and Tmall have focused on supporting "original-design" merchants, including brand merchants, while unlocking the potential of high-quality white-label supply from key industrial clusters.

We will continue to strengthen our partnerships with leading brands, helping them achieve stable and sustainable business growth. Tmall remains the primary platform on which brands and many "original-design" merchants operate. At the same time, we will dive deeper into industrial clusters to source high-quality products, and support more manufacturing factories to operate directly on our platform.

We will also leverage our platform's AI capabilities to enable more white-label merchants to adopt a simpler, more efficient managed operation model. The managed operation model for industrial-cluster merchants continues to account for a growing share of platform transactions. During the recent 618 shopping festival, despite certain macroeconomic challenges, the results were in line with our expectations. Notably, core merchants still achieved quite good results.

At the same time, we see opportunities to apply AI across both the supply and demand sides of e-commerce. On the consumer side, we will continue to launch new AI-powered experiences and scenarios, such as multimodal search and AI virtual try-ons. We will also use AI to enhance the experience and efficiency of existing shopping scenarios. For example, we have already seen AI deliver significant improvements in product recommendations.

On the merchant side, AI is already widely used in merchants' operations. We are applying AI across different areas of their operations to strengthen their capabilities, particularly in data analytics, advertising and marketing, and customer service, where merchants can derive clear benefits. Going forward, we will also collaborate with QwenWork to launch AI agents tailored for e-commerce use cases.

Turning to quick commerce: after more than a year of investment and development, Taobao Instant Commerce has seen material changes in both scale and market share, with significant improvements in user mindshare, supply diversity, logistics experience, and order volume. Last quarter (June Quarter), while maintaining growth in both users and order volume, we substantially improved UE, and significantly narrowed losses.

Building on this foundation, we will move faster to integrate related businesses such as Freshippo and Tmall Supermarket and expand non-food quick commerce, with a particular focus on scaling our front warehouse network. Over the past year, Freshippo has stepped up its front warehouse rollout, helping accelerate year-over-year GMV growth across the business.

Meanwhile, our quick commerce will continue to expand its category coverage and innovate in key categories to enhance the consumer experience. We expect quick commerce transaction volume from non-food categories to surpass food categories within the next fiscal year, driving growth in many physical goods categories across the overall e-commerce platform. The quick commerce segment is expected to achieve overall profitability in FY29. In the long term, we believe quick commerce has the potential to contribute 30% of the platform's total GMV, becoming the second growth curve for our e-commerce business.

On International E-commerce: in the short term, our international e-commerce business has indeed been affected by international tariff policies and the geopolitical environment, and is facing some pressure. Nevertheless, despite the complex market environment, our cross-border e-commerce business has continued to grow its transaction volume while significantly improving profitability. In terms of both transaction volume and profitability, we believe the cross-border business has long-term growth potential. In addition, our e-commerce platforms in regions such as Türkiye and the Middle East are growing rapidly, and operating efficiency in markets such as Southeast Asia continues to improve.

Finally, on Global Wholesale. Our B2B businesses, including our 1688 platform and Alibaba.com, have sustained continued growth over the past two decades. We believe that AI technology will bring profound changes to B2B platforms, and may even fundamentally change their existing business models. In particular, AI agents will play an increasingly important role in B2B transactions. Accio Work, the AI agent we launched for cross-border merchants, attracted 50,000 paying users shortly after its launch.

AI is reshaping how B2B merchants — especially cross-border merchants — do business. We believe that, building on the capabilities and know-how we have built up over more than 20 years in this field, we have the opportunity to create entirely new business models and commercial opportunities in B2B and cross-border trade in the AI era.

Overall, over the past few years, we have completed a phase of strategic build-out across several areas of our e-commerce business. Going forward, we aim to continue leveraging our strengths — from supply-chain synergies to AI technology — to unlock greater potential across these businesses in the AI era, while developing a more diversified revenue and profit mix and driving steadier growth across our e-commerce business as a whole.

Yang Bai (CICC):

(Translation)

My question concerns the cloud and AI business. We have observed that Alibaba Cloud's revenue growth has continued to accelerate, reaching 45% this quarter. The company

previously set out a long-term goal of surpassing US$100 billion in external Cloud revenue within the next five years, and today you also indicated that growth will continue to accelerate over the next few quarters.

I have two questions. First, looking at the growth trajectory of the cloud business over the coming quarters, what are the core drivers supporting a further acceleration in cloud computing growth? Second, the industry is still in a phase of tight compute supply. Management just mentioned that the supply-demand picture may not change before 2030. But taking a longer-term view, what are the core drivers of the cloud business's long-term growth, and do they differ from the short-term ones?

Eddie Wu:

(Translation)

Let me elaborate on this a bit, in three parts: first, our current position and key metrics; second, the growth drivers for our AI-related business; and third, our longer-term outlook.

First is our current position and key metrics.

Growth in external revenue from our AI and cloud business has accelerated for nine consecutive quarters, reaching 45% this quarter. Customer demand is strong, and our supply capabilities give us a significant advantage over other cloud providers, so we expect revenue growth to continue to accelerate over the next few quarters. AI-related products generated RMB12.4 billion in revenue this quarter, equivalent to an annualized revenue run rate of approximately US$7.3 billion. Based on our current forecasts, the annualized revenue run rate of our AI products will reach US$10 billion next quarter. So you can see that growth remains very strong.

At the same time, we expect the EBITA margin of our Cloud business to improve steadily each quarter over the next few quarters. As for MaaS, which is an important part of our Cloud business, growth in AI demand and improvements in inference efficiency drove ARR to surpass RMB16 billion as of August.

Now, the growth drivers based on our business model.

The biggest difference between Alibaba's approach to AI investment and that of pure-play AI companies is that we invest at scale across the full AI technology stack, particularly in chips, AI cloud infrastructure, and models. That has secured us an industry-leading position across all three of these most critical layers.

In our view, the AI industry is still at an early stage of technological development, and at different stages of that development, the industry's core commercial value will shift among different layers — chips, cloud, models, and applications. Our full-stack investment ensures that we can deliver leading service capabilities and price-performance, enabling us to remain competitive and sustain long-term growth momentum through every stage of the industry's future development.

Let me now turn to what we see as the most important growth driver in the near term, in particular over the next year or two.

The large-scale commercialization of inference services, which took off at the end of 2025, has fundamentally changed the business model of AI cloud. As I mentioned earlier, compute is the core asset underpinning AI monetization. Every AI revenue model today is built around AI compute, and the industry now broadly expects compute capacity to remain constrained over the medium term. At the same time, given the high margins of MaaS inference services, AI compute is no longer a traditional cost center but a productive asset directly linked to revenue generation.

High-priced compute capacity remains in short supply across the industry, while GPUs are required across a wide range of use cases. Against this backdrop, pricing across these use cases is increasingly being aligned with the industry's highest-margin forms of monetization, driving a broader shift in pricing models across nearly all GPU-related products.

Alibaba possesses model capabilities spanning all modalities, most at the industry frontier. This strengthens our ability to monetize compute and provides a strong pricing anchor, giving Alibaba Cloud differentiated pricing power for compute. With demand remaining robust, the market is repricing the commercial value of compute. We are able to sign new contracts at healthier pricing, and renewals of existing customer contracts are likewise being repriced in response to new market demand. As a result, Alibaba Cloud's profitability has been improving significantly since the beginning of this year.

The economies of scale and network effects of the AI cloud business model are important long-term growth drivers. Over the past two years, investors have kept asking what the AI super application will be. We believe the biggest AI super application is an AI cloud platform centered on AI compute.

In the long run, most enterprises will rely on full-stack AI cloud platforms to support their core AI workloads, including large-scale training, large-scale inference, customized inference software, and the development and operation of enterprise AI agents. These workloads require a full-stack cloud platform providing comprehensive cloud-based infrastructure services, including GPUs, CPUs, storage, databases, virtualization, and harness tools.

An AI cloud platform is like a supercity: training and inference workloads, enterprise systems, and AI agents are its residents, while full-stack AI cloud services form its infrastructure. Continually upgrading and improving that infrastructure helps attract a constant stream of new residents and retain existing ones.

Over the long term, the economies of scale and network effects of our AI cloud platform will be enduring drivers of growth for our AI business.

We operate the largest number of data centers among Asian cloud providers, giving us the strongest economies of scale. As we deploy our proprietary T-Head AI chips at greater scale, we can avoid the margin erosion caused by the high premiums on externally sourced, high-end

GPUs. And the SOTA capabilities of our proprietary models give us very strong pricing power for compute.

Based on both industry trends and our own product advantages, we see very strong long-term momentum for revenue growth and margin expansion in AI and cloud. We are therefore highly confident in our ability to accelerate progress toward our target of US$100 billion in external cloud revenue by calendar year 2030, and we also have high visibility into achieving a profit margin of over 20% in our AI cloud business.

Yuan Liao (CITIC Securities):

(Translation)

Congratulations on the positive progress in cloud and AI this quarter.

I have a follow-up question on the MaaS business. As Eddie mentioned just now, MaaS ARR exceeded RMB16 billion in August. Last quarter, the company set a target of surpassing RMB30 billion in MaaS ARR by year-end. Given the progress to date, will you adjust year-end target?

In addition, within the MaaS business, what are the respective shares of Qwen models versus third-party models? Looking ahead, how will intensifying model competition and the emergence of more open-weight models affect MaaS gross margins and profitability?

Eddie Wu:

(Translation)

Our MaaS business on Alibaba Cloud's Model Studio is indeed growing very rapidly. MaaS ARR surpassed RMB16 billion in August. Given the current growth momentum and the new models we will continue to launch, we have clear visibility towards the target of surpassing RMB30 billion in ARR by year-end.

In our MaaS business today, proprietary models still account for the majority, although revenue from third-party open-weight models is also substantial.

Let me also share our longer-term view on model competition.

There is strong customer demand for using models from multiple providers within a single AI application, as each model has its own distinctive capabilities. At the same time, we also expect more open-weight models to emerge. This trend benefits inference platforms offered by cloud providers, such as Alibaba Cloud’s Model Studio, as gross margins are broadly comparable whether we host freely available third-party open-weight models or our own proprietary models.

We develop our proprietary models primarily to advance model intelligence over the long term and pursue breakthroughs toward AGI. For MaaS services in the near term, proprietary and

hosted open-weight models have similar gross-margin profiles. Therefore, a more competitive, or a more open and vibrant model ecosystem is actually beneficial to cloud providers such as Alibaba Cloud.

Alex Yao (JP Morgan):

(Translation)

Eddie just spent quite some time discussing the full-stack AI ecosystem. My question is: at which layer of this full-stack AI ecosystem will value ultimately accrue or be monetized?

We saw that the company opened the weights of Qwen3.8-Max within three months of the flagship model's release, while the commercialization of your proprietary chips has also been quite successful, with the chips now serving more than 650 external customers. Does this mean management believes value will ultimately accrue to the compute and scheduling and orchestration layer, rather than the model layer?

Or does management believe that at different stages of AI development, the value of the full-stack ecosystem may accrue to different layers? If you believe long-term value will accrue to the hardware or compute layer, how do you think about protecting returns at that layer once government-led compute supply expands, or competitors that also have compliant chips expand their supply?

Eddie Wu:

(Translation)

This involves making a long-term judgment. On that judgment, I think some uncertainty still remains. But overall, because we invest across the full AI stack, we believe that we have the opportunity to capture value within our ecosystem regardless of which layer it accrues to, or how that distribution shifts across different layers at different stages.

Let me share my own near-term assessment. For now, we believe most of the value remains concentrated in the chips and cloud infrastructure layer. We see the same pattern across many companies both in China and overseas: when a technology is in its early stages, especially when its supply is scarce, a large share of the value typically accrues to providers of infrastructure and critical hardware such as chips and storage. Under our new segment structure, compute, cloud infrastructure, and AI inference are grouped together within a single segment.

As for your question on how commercial value will be realized at the model layer, there is considerable debate within the industry, and views also differ within our company.

In my view, API-based monetization of foundation models is a short-term, transitional business model, not the ultimate one. We are not making such large investments in compute and full-

stack model development simply to generate near-term API revenue. We believe that once foundation models achieve or approach AGI, the business model will shift toward directly creating products and directly delivering the outcomes customers need, particularly in product R&D, product creation, and business operations. At that stage, AI model companies should no longer center their business models on providing API services. It is the prospect of that much greater commercial value in the future that is driving so many model companies around the world to continue investing in today's "arms race". From this perspective, I believe the commercial value of foundation models is still far from being realized. Their future commercial value will be far greater than that of the API services provided today.

Turning to the hardware and compute layer, I would like to address your question in some depth. Previously, we have not elaborated much during earnings calls about our proprietary chip strategy and T-Head, so let me elaborate a bit today.

T-Head has shipped more than 500,000 units of its previous-generation chips. Separately, the first generation of our chips backed entirely by domestic supply chain, which I mentioned earlier, began a broad commercial rollout on Alibaba Cloud in August through supernode instances. I believe that makes us one of only two companies in the Chinese market capable of commercially offering supernode instances at scale. We are highly confident in this generation of domestic chips, both in terms of customer adoption and large-scale commercialization. I should also highlight one feature that distinguishes T-Head among domestic chips. To my knowledge, T-Head's AI chips are unique in adopting GPGPU as their core architecture, enabling them to support both training and inference workloads effectively. We currently serve several hundred mid-to-large customers across both training and inference workloads, including embodied AI companies, autonomous driving companies, and large model companies. As such, we may be the only company in China with commercial customers running both training and inference at scale.

In addition, I would like to share our expectations for T-Head's second-generation domestic chip.

In the second half of this year, T-head's second-generation domestic chip is gradually taping out with production to follow at a later stage. The chip will offer very strong compute performance and interconnect bandwidth, and we believe it will be fully capable of supporting large-scale model training. This gives us a truly unique position in domestic chip technology. I also do not think there is a so-called government-led compute supply market. We remain very confident that T-Head will become a core competitive advantage for Alibaba Cloud, especially given that several hundred mid-to-large customers are already realizing meaningful value from using T-Head chips through Alibaba Cloud. As I know, our chips also have broad reach across China's engineering community.

To sum up, I would rank T-Head first among domestic chips in terms of the breadth and diversity of training and inference workloads they can support across industries. In terms of our build-out of future production capacity and supply chains, we are among the top two players in the industry. And Alibaba Cloud's longstanding market-share leadership in China's AI cloud market gives us very strong channels for bringing AI chips to customers in the future. We therefore remain highly confident in the long-term commercial value of T-Head chips.

Lydia Liu:

(Original)

Thank you very much. We appreciate your support, and we look forward to updating you on our progress next quarter. Thank you.

[END OF TRANSCRIPT]